

SPEECH LAB EVAL

NAME : Dev Jalota

ROLL NO. : 102103452

SUBGROUP : 4CO17

Paper Summary :

The paper "**Speech Commands: A Dataset for Fine-Grained Speech Recognition**" presents a dataset with 65,000 one-second recordings of 12 spoken commands, such as "yes," "no," "up," and "down." It aims to advance research in speech recognition by offering a standardized benchmark for fine-grained audio classification, focusing on real-world performance in command recognition.

Google Collab files link :

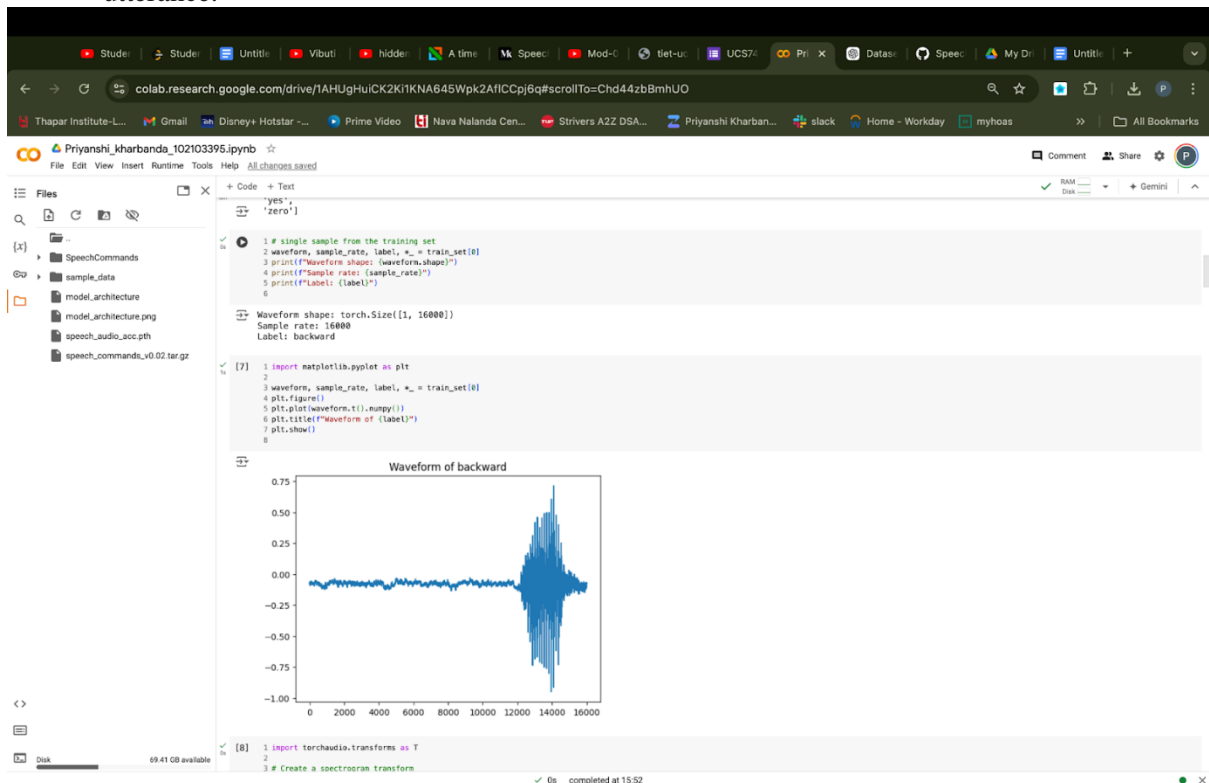
<https://colab.research.google.com/drive/1k2v-4TAplImoT-EbDsSzEr4eQBCU640#scrollTo=SdDDteRj07ft>

https://colab.research.google.com/drive/1dcbddNXmg2VLQ_4t2HFQZc86bKP_z6cC

https://colab.research.google.com/drive/1dcbddNXmg2VLQ_4t2HFQZc86bKP_z6cC

Analysis of data

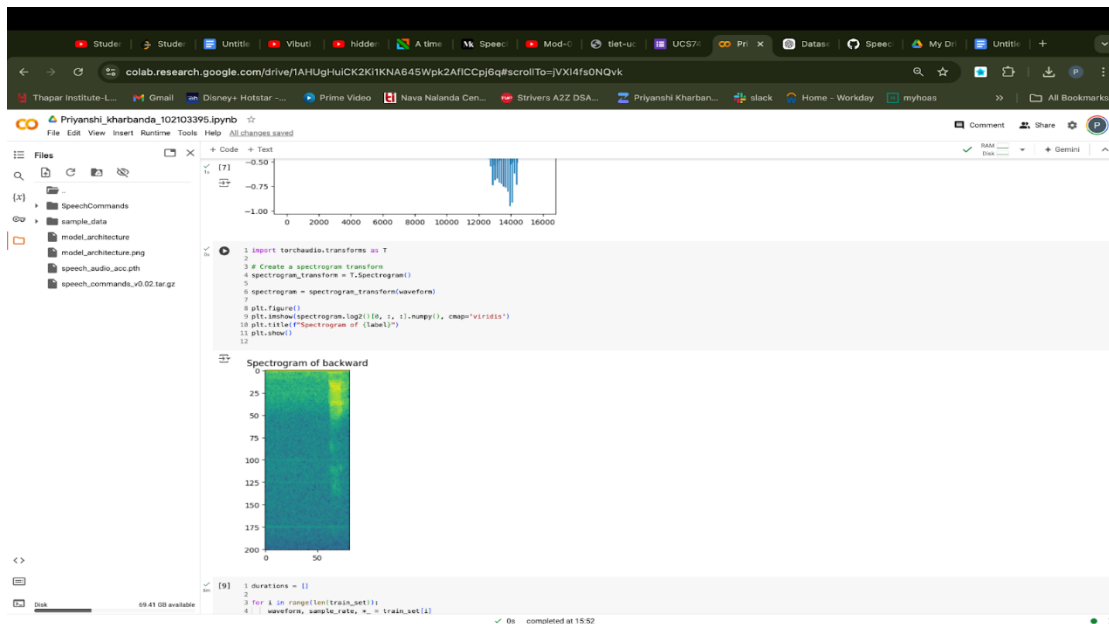
1. A data point in the SPEECHCOMMANDS dataset is a tuple made of a waveform (the audio signal), the sample rate, the utterance (label), the ID of the speaker, the number of the utterance.



Number of training samples: 84843

Number of testing samples: 11005

2. After applying fourier transformation, here is what a spectrogram of label backward looks like



3. Analysis done of duration of samples in our speechcommand dataset.

```

1 durations = []
2
3 for i in range(len(train_set)):
4     waveform, sample_rate, *_ = train_set[i]
5     durations.append(waveform.shape[1] / sample_rate)
6
7 print(f"Average duration: {sum(durations) / len(durations)} seconds")
8 print(f"Maximum duration: {max(durations)} seconds")
9 print(f"Minimum duration: {min(durations)} seconds")
10

```

Average duration: 0.9807155429734978 seconds
Maximum duration: 1.0 seconds
Minimum duration: 0.256 seconds

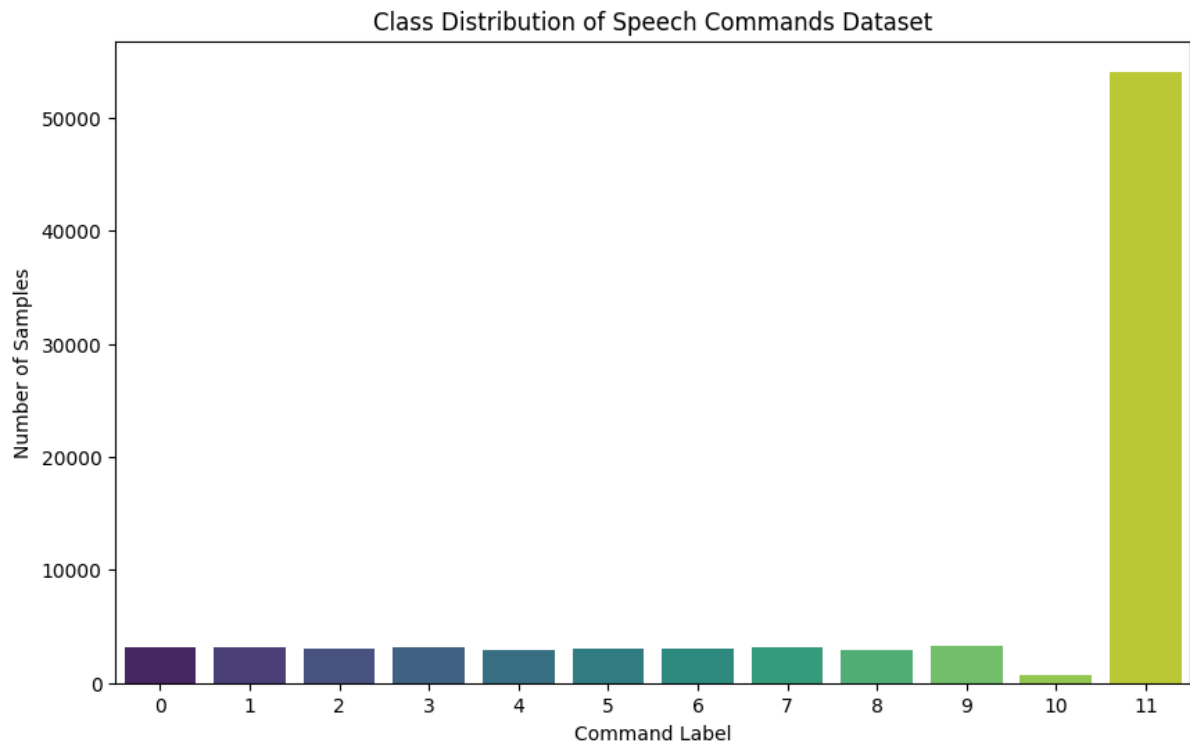
```

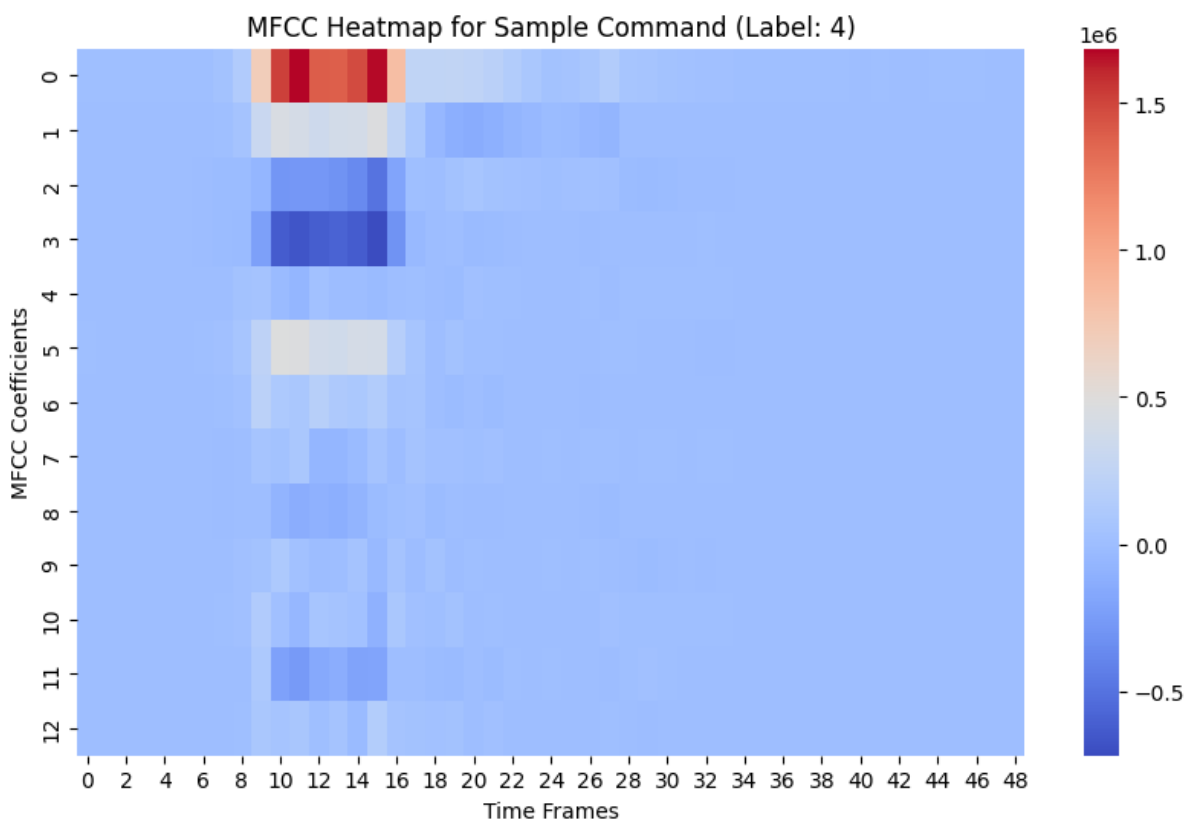
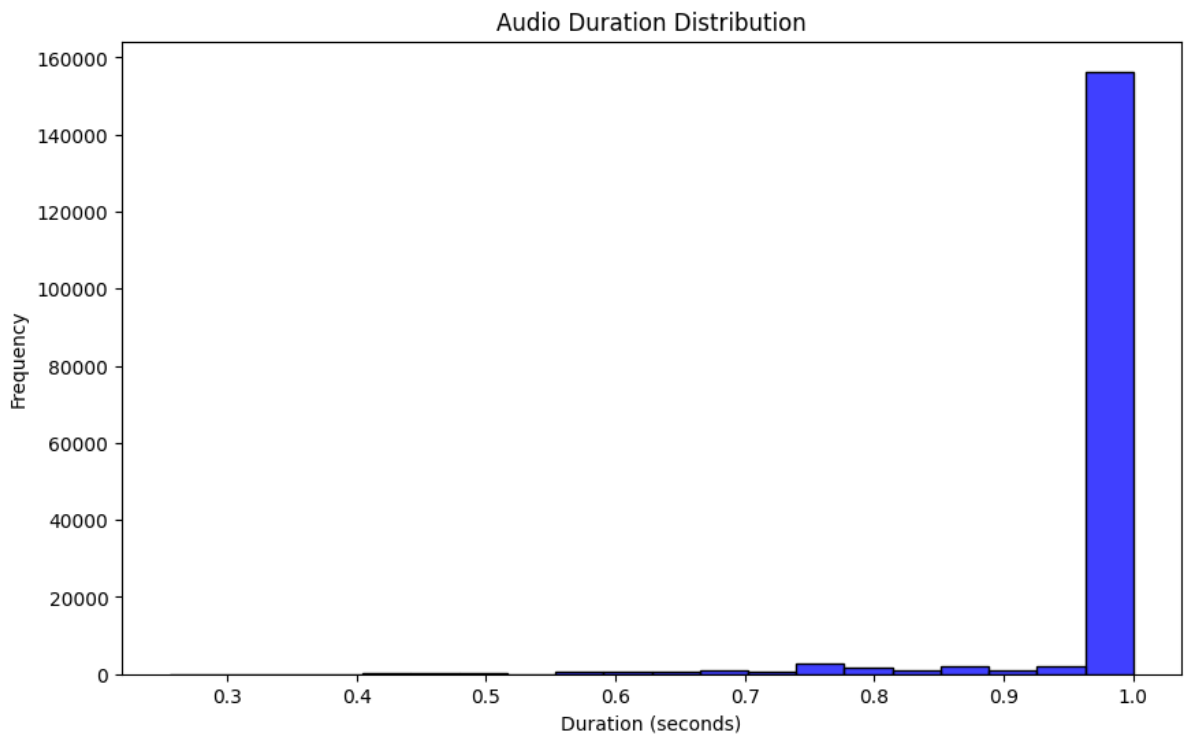
[10] 1 waveform_first, *_ = train_set[0]
2 ipd.Audio(waveform_first.numpy(), rate=sample_rate)
3
4 waveform_second, *_ = train_set[1]
5 ipd.Audio(waveform_second.numpy(), rate=sample_rate)

```

0:00 / 0:01

Some snippets of analysis :





Steps considered while setting up data loading and batching for training and testing a model using the **SPEECHCOMMANDS** dataset with PyTorch.

2. Padding Sequences (**pad_sequence** function)

- **Purpose:** Ensures all tensors in a batch are the same length by padding with zeros.
- **Details:**
 - Transposes each tensor in the batch.
 - Uses **torch.nn.utils.rnn.pad_sequence** to pad tensors to the same length.
 - Permutes the dimensions of the padded tensor to match the expected input shape for your model (usually **(batch_size, channels, sequence_length)**).

3. Collate Function (**collate_fn** function)

- **Purpose:** Defines how to collate (combine) individual data samples into a batch.
- **Details:**
 - Extracts waveforms and labels from each sample in the batch.
 - Applies **pad_sequence** to ensure waveforms in the batch have the same length.
 - Converts labels to indices using **label_to_index** and stacks them into a tensor.

4. DataLoader Configuration

- **Purpose:** Creates **DataLoader** instances for the training and test datasets.
- **Details:**
 - **train_loader:** Loads training data with batching, shuffling, and custom collate function.
 - **batch_size=256**
 - **shuffle=True:** Shuffles data at each epoch.
 - **num_workers** and **pin_memory** are set based on whether a GPU (**cuda**) or CPU is used.
 - **test_loader:** Loads test data with batching and custom collate function.
 - **batch_size=256**
 - **shuffle=False:** No shuffling for test data.
 - **drop_last=False:** Keeps the last batch even if it's smaller than the batch size.

Model Used (**M5** Class)

Purpose: Defines a CNN model for classification.

Components:

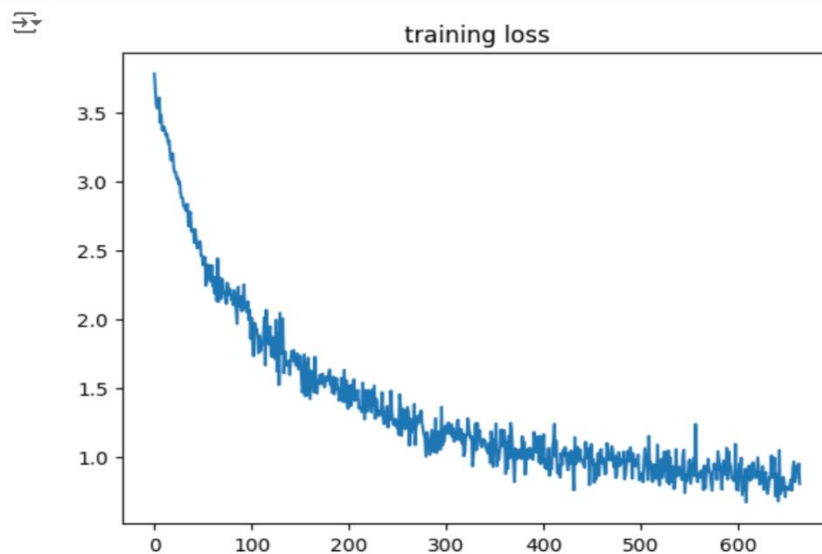
- **Convolutional Layers:**
 - **conv1:** 1D convolution with kernel size 80, stride 16.
 - **conv2:** 1D convolution with kernel size 3.
 - **conv3:** 1D convolution with kernel size 3, output channels doubled.
 - **conv4:** 1D convolution with kernel size 3, output channels doubled.
- **Batch Normalization:**
 - **bn1, bn2, bn3, bn4:** Normalize outputs of respective convolutional layers.
- **Pooling Layers:**
 - **pool1, pool2, pool3, pool4:** Max pooling with kernel size 4.
- **Fully Connected Layer:**
 - **fc1:** Linear layer that maps from the output channels to the number of classes.

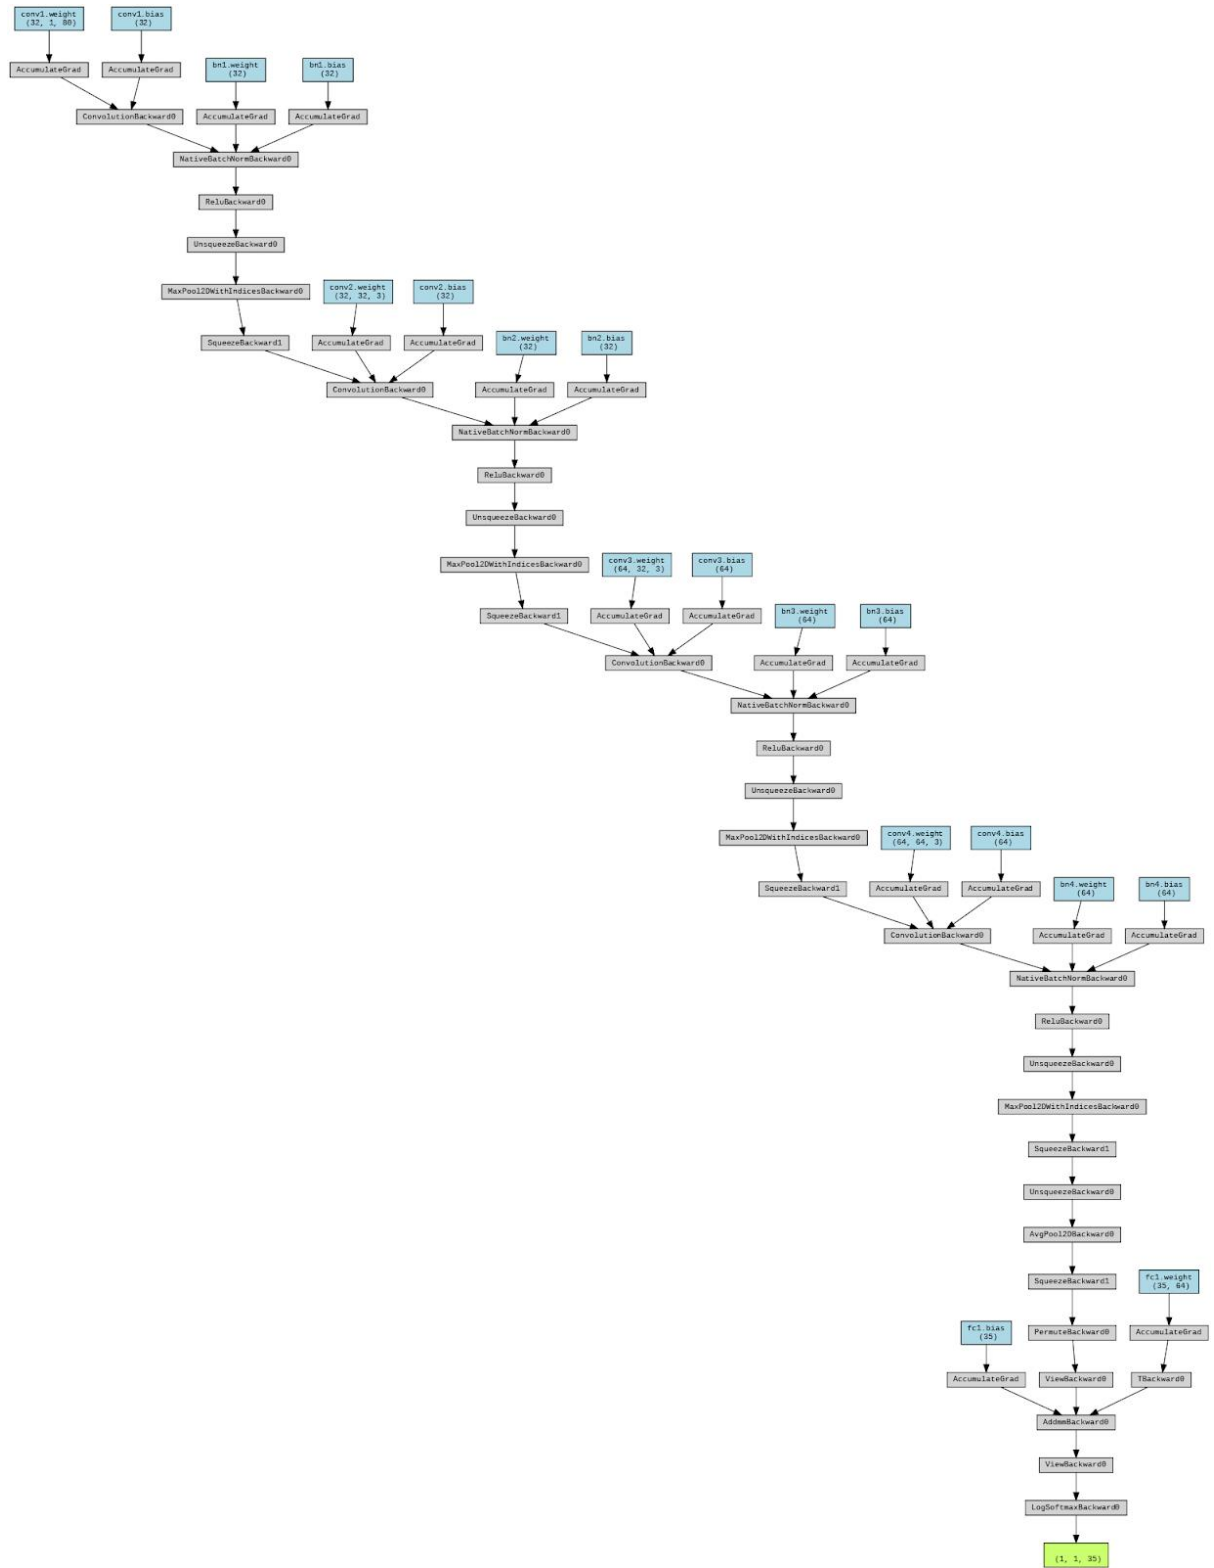
- **Forward Pass:**

- Applies convolutions, batch normalization, ReLU activation, pooling, and average pooling.
- Reshapes and passes through a fully connected layer followed by `log_softmax` for classification.

```
90% |██████████| 1.80266666666666513/2 [17:05<01:44, 528.12s/it]Trai
93% |██████████| 1.85599999999999832/2 [17:35<01:13, 510.27s/it]Trai
100% |██████████| 1.99999999999999793/2 [18:54<00:00, 567.26s/it]
Test Epoch: 2 Accuracy: 5745/11005 (52%)
```

```
✓ 0s 1 plt.plot(losses);
2 plt.title("training loss");
```





Model M5 Architecture

Results – 72.53 % Accuracy